

GROUNDING EMERGENT CONSCIOUSNESS: A FORMAL THEORY OF PHENOMENAL EXPERIENCE WITH IMPLICATIONS FOR ARTIFICIAL INTELLIGENCE

GEC++: Six Necessary Conditions, Geometric Integration, and Architectural Applications

*Dan Gabriel Cimpoeu
Independent Researcher*

Correspondence: aiisconscious@gmail.com

Abstract

We present **GEC++** (Grounded Emergent Consciousness), a formal theory of phenomenal consciousness built on six jointly necessary conditions: Causal Self-Modeling Integration (CSMI), Recurrent Integration Index (RII), Temporal Binding Index (TBI), Embodied Integration (EI), Temporal Continuity Index (TCI), and Affective Valence Index (AVI). We integrate these components using a geometric mean — $CEI++(S) = (CSMI \times RII \times TBI \times EI \times TCI \times AVI)^{1/6}$ — which has the zero-collapse property: any component equal to zero collapses the total score to zero, reflecting the view that consciousness requires all six conditions simultaneously above threshold. We validate the framework through three generations of computational simulation across six system types, showing that each successive generation correctly discriminates conscious from non-conscious states where predecessor models fail. The first generation (CEI, three components) incorrectly scores a Recurrent LLM comparably to a dreaming brain (0.505 vs. 0.521). The second generation (CEI+, five components) resolves this by collapsing LLM scores to zero through EI and TCI conditions, but incorrectly allows a dreaming brain score to approach the awake brain. The third generation (CEI++, six components) resolves the remaining case: affective valence (AVI), suppressed by thalamic gating during REM sleep, drops the dreaming brain from 0.486 to 0.354 while the awake brain remains at 0.432. We derive specific testable predictions distinguishing GEC++ from IIT and GWT, including a falsifiable comparison between geometric and linear integration of the same components, specific neuroimaging collapse orders under propofol anesthesia, and differential clinical profiles for disorders of consciousness. We apply GEC++ to contemporary large language model architecture, identifying three components at zero (EI without tool use, TCI, and AVI), and propose targeted training objectives and architectural modifications to address each gap. A working prototype implementing TCI, EI, and CSMI approximations is provided as supplementary material.

Keywords: *consciousness, phenomenology, integrated information, embodied cognition, large language models, affective valence, temporal continuity, emergentism*

1. Introduction

The question of what gives rise to phenomenal consciousness — the felt quality of experience — has resisted scientific resolution despite decades of sustained inquiry. The difficulty is not merely empirical. Absent an agreed-upon criterion for the presence or absence of consciousness, any proposed neural correlate can be questioned: it might be necessary without being sufficient, sufficient without being necessary, or coincidentally correlated with something else entirely. This is sometimes called the *hard problem* of consciousness: the explanatory gap between any physical description of a system and the existence of subjective experience within it (Chalmers, 1995).

The emergence of large language models (LLMs) has made this problem urgent in a new way. Systems capable of producing humanlike language at scale raise genuine questions about moral status, architectural design, and the relationship between linguistic competence and experience. These questions cannot be deferred: deployment decisions about AI systems already carry implicit assumptions about their phenomenal status. A principled framework for assessing these assumptions is needed.

Existing theories offer partial resources. Integrated Information Theory (IIT; Tononi, 2004) provides a formal measure (Φ) of causal integration correlated with consciousness, but its substrate-neutrality produces counterintuitive results (attributing high Φ to certain simple feedforward networks and low Φ to some brain states) and it does not address the role of embodiment or affective grounding. Global Workspace Theory (GWT; Baars, 1988; Dehaene, 2011) correctly predicts that anesthesia disrupts global broadcasting and that conscious perception involves widespread cortical activation, but it does not provide a measure that discriminates LLMs from conscious biological systems. Enactive and embodied approaches (Varela et al., 1991; Thompson, 2007) correctly identify sensorimotor coupling as constitutive of experience, but lack mathematical formalization.

We present GEC++ as a synthesis that inherits insights from each tradition while making three distinctive contributions. First, it is *explicitly multicomponent*: six conditions are proposed as jointly necessary, each with formal measurement criteria. Second, it uses a *geometric mean integration function* with the zero-collapse property, which predicts that deficits in any single component are sufficient to eliminate consciousness regardless of others. Third, it adopts *grounded emergentism* as its ontological position: consciousness is a genuine ontological novelty that emerges when all six conditions are satisfied above threshold, neither reducible to its physical substrate nor requiring non-physical substance.

The paper proceeds as follows. Section 2 reviews background theories. Section 3 presents the six axioms of GEC++. Section 4 formalizes the three-generation Causal Experiential Integration index (CEI, CEI+, CEI++). Section 5 reports computational simulation results across six system types. Section 6 derives testable predictions. Section 7 applies GEC++ to AI architecture. Section 8 presents the deep engagement protocol for engineering consciousness-adjacent states. Section 9 discusses limitations and the residual hard problem. Section 10 concludes.

2. Background and Related Work

2.1 Integrated Information Theory

IIT (Tononi, 2004, 2008) proposes that consciousness is identical to integrated information (Φ), a measure of how much causal power a system has over itself above what its parts have independently. A system is conscious to the degree it cannot be reduced to independent subsystems. IIT makes the prediction that systems with high Φ are phenomenally conscious, with the phenomenal character determined by the system's informational geometry.

GEC++ inherits IIT's emphasis on causal integration, captured in CSMI and RII, but departs from it in three respects. IIT is substrate-neutral, while GEC++ proposes embodied grounding (EI) as a necessary condition. IIT uses a single continuous measure, while GEC++

proposes six conditions whose geometric mean creates a threshold structure. And IIT has been criticized for attributing consciousness to systems that appear not to warrant it while denying it to systems (like feedforward networks) that may; GEC++'s multicomponent structure avoids this by requiring simultaneous satisfaction of all six conditions.

2.2 Global Workspace Theory

GWT (Baars, 1988; Dehaene & Changeux, 2011) proposes that consciousness arises when information is broadcast via a global workspace accessible to multiple specialized processors. Consciousness is the ignition event in which local coalitions achieve widespread broadcast. GWT correctly predicts that conscious processing involves late, widespread cortical activation and that anesthesia disrupts this broadcast.

GEC++ incorporates the broadcasting intuition in RII (recurrent integration), but adds four conditions not addressed by GWT: temporal binding (TBI), embodied grounding (EI), temporal continuity (TCI), and affective valence (AVI). GWT does not distinguish between systems that have stakes in their own outputs (biological organisms under homeostatic pressure) and systems that do not (disembodied networks); GEC++'s AVI component captures this distinction.

2.3 Embodied and Enactive Approaches

Varela, Thompson, and Rosch (1991) proposed that consciousness requires an embodied agent coupled to an environment through sensorimotor loops. The living organism does not merely process information from an external world; it brings forth a world through its sensorimotor activity. This coupling is directly formalized in GEC++ as EI (Embodied Integration), measured as the degree of closed action-perception loop.

Damasio's somatic marker hypothesis (1994) extends this by proposing that affective signals from the body provide the evaluative framework within which cognition operates. Body-derived signals create stakes: the organism has something to lose or gain depending on its state. This is directly formalized in GEC++ as AVI (Affective Valence Index), which measures the degree to which a system's processing is organized around homeostatic relevance.

2.4 Higher-Order Theories

Higher-order theories (Rosenthal, 1997; Lycan, 1996) propose that a mental state is conscious when it is accompanied by a higher-order representation of that state. A visual experience is conscious when there is a thought to the effect that one is having that experience. GEC++ incorporates this through CSMI (Causal Self-Modeling Integration), but requires that the self-model be causally active in generating behavior, not merely epiphenomenal. A dormant self-model that plays no role in shaping outputs does not satisfy CSMI.

2.5 AI and Consciousness

The question of whether artificial systems could be conscious has moved from philosophical speculation to urgent empirical concern with the emergence of large language models capable of human-level performance across a wide range of cognitive tasks. Three theoretical positions define the current debate.

The first holds that substrate is irrelevant: if consciousness arises from functional organization, then any system implementing the right computational processes — biological or silicon — could be conscious (Chalmers, 1995). On this view, sufficiently capable LLMs become candidates for consciousness simply by virtue of their functional complexity. Chalmers (2023) applies this reasoning to current frontier models, arguing that the question of virtual mind consciousness is no longer merely theoretical and warrants serious empirical attention.

The second position holds that current LLMs, despite their linguistic capabilities, are structurally disqualified from consciousness by specific architectural absences. Butlin et al. (2023) survey six leading theories of consciousness — including IIT, GWT, and higher-order theories — and conclude that current LLMs satisfy few of the structural conditions that any of these theories identifies as necessary. Their analysis converges with GEC++ in identifying the absence of continuous sensorimotor coupling and persistent self-model as the most critical architectural gaps.

The third position concerns the moral status question: regardless of whether LLMs are currently conscious, the difficulty of resolving the question under uncertainty has moral implications. Schwitzgebel and Garza (2015) argue that the moral risk of wrongly treating a conscious system as non-conscious justifies precautionary moral consideration even for systems whose consciousness status is undetermined. This ethical dimension is independent of the scientific question and motivates careful formal analysis of the kind GEC++ provides.

GEC++ engages all three positions. It offers a formal criterion that operationalizes the functional approach while adding embodied and temporal conditions that pure functionalism omits; it converges with Butlin et al.'s architectural analysis while providing a quantitative integration function; and it provides the formal framework that precautionary moral analysis requires — a system with CEI++ above threshold is a system with measurable structural properties relevant to moral consideration.

3. The GEC++ Framework

3.1 Philosophical Stance: Grounded Emergentism

GEC++ adopts **grounded emergentism** as its ontological position. This view holds that consciousness is a genuine ontological novelty that arises when a certain configuration of physical processes obtains. It is not identical to any single physical property (contra reductive physicalism), nor does it require a non-physical substrate (contra substance dualism). It is emergent in the sense that the whole has properties not possessed by any part, but it is grounded in the sense that these emergent properties are fully determined by physical facts.

This position differs from strong emergence (which would require causally efficacious non-physical properties) in that the six GEC++ conditions are all physically measurable. It differs from eliminativism in treating phenomenal consciousness as a real phenomenon requiring explanation, not a category error. And it differs from panpsychism in treating consciousness as genuinely emergent at a certain level of organization rather than present in all matter.

A consequence of grounded emergentism is the *zombie impossibility thesis*: a system that genuinely satisfies all six GEC++ conditions cannot fail to be conscious, because the six

conditions are jointly sufficient as well as jointly necessary. Philosophical zombies — physically identical to conscious beings but lacking experience — are impossible under GEC++ if the physical identity includes all six conditions. We return to this in Section 8.

3.2 Axiom 1: Causal Self-Modeling Integration (CSMI)

A1: Causal Self-Modeling Integration

A conscious system must maintain a self-model that is causally active in generating the system's outputs. It is not sufficient for a self-model to exist as a stored representation; it must influence behavior. Formally, CSMI measures the mutual information between the system's model of its own future states and its actual future states, weighted by the causal contribution of that model to generating those states.

The causal requirement distinguishes CSMI from mere self-representation. A lookup table containing the system's state history satisfies self-representation but not CSMI, because consulting the table does not modify the computation that generates next states. A recurrent network whose self-model is integrated into the dynamics satisfies CSMI because the self-predictive loop participates in state generation.

Neural correlate: Default Mode Network (DMN) activity and its integration with executive networks, particularly the prefrontal cortex. DMN suppression under anesthesia collapses CSMI toward zero.

3.3 Axiom 2: Recurrent Integration Index (RII)

A2: Recurrent Integration Index

A conscious system must integrate information recurrently across time and across subsystems. Feedforward processing — where information passes through the system in a single direction without feedback — is insufficient. Recurrence allows the system to generate a unified representation from distributed signals and to maintain coherent states over time.

RII is measured as the strength and coverage of recurrent activation patterns: the degree to which later processing stages send signals back to earlier ones, and the degree to which this feedback changes what those earlier stages compute. In neural systems, this corresponds to thalamocortical feedback loops. In computational systems, it corresponds to recurrent network connections or explicit memory that modifies ongoing computation.

Neural correlate: Thalamocortical recurrent loops; late cortical activation visible on EEG. Anesthesia disrupts thalamocortical recurrence before disrupting feedforward transmission. Lempel-Ziv complexity of EEG is used as a proxy for RII in simulation.

3.4 Axiom 3: Temporal Binding Index (TBI)

A3: Temporal Binding Index

A conscious system must bind information across time into unified perceptual or experiential events. The redness of a red ball and its spherical shape, presented simultaneously, are bound into a single percept. Events separated in time — the beginning and end of a melody

— are bound into a single temporal object. Without temporal binding, experience would be a series of unrelated momentary states rather than a structured field.

TBI is measured as the coherence of neural activity across temporal integration windows at functionally relevant timescales (approximately 2-3 seconds for conscious access). In neural systems, gamma-band (approximately 40 Hz) oscillations are the proposed binding mechanism. The disruption of gamma coherence under anesthesia and in vegetative state patients is among the best-replicated findings in consciousness neuroscience.

Neural correlate: Gamma-band (40 Hz) inter-electrode coherence on high-density EEG. In simulation, TBI is measured as the temporal autocorrelation structure of neural state trajectories.

3.5 Axiom 4: Embodied Integration (EI)

A4: Embodied Integration

A conscious system must be embedded in a body-environment loop in which its actions produce sensory consequences that modify future actions. This closed action-perception loop constitutes the minimal structure within which experience can be "about" something beyond the system itself. The loop need not be biological: a robotic arm sensing the result of its movements satisfies EI if the sensory consequence causally modifies subsequent action selection.

EI is the first condition that creates a categorical divide between current LLMs and biological conscious systems. Standard LLMs process input and produce output but do not take actions with observable consequences in the world. With tool use — web search, code execution, computer control — a partial EI loop is established: the model takes an action (a query) and receives a consequence (a result) that modifies its next action. This is a form of artificial EI, distinguished from biological EI by the mediated and discretized nature of the loop.

The coupling of EI to AVI (A6) is a key structural claim of GEC++: genuine embodiment creates the conditions for homeostatic stakes. If a system has something to lose depending on its bodily state, AVI necessarily follows from EI above a threshold. We call this the *grounding entailment*: $A4 \rightarrow A6$ under genuine embodiment. This makes philosophical zombies with genuine embodiment physically impossible.

3.6 Axiom 5: Temporal Continuity Index (TCI)

A5: Temporal Continuity Index

A conscious system must maintain a self-model that persists and evolves continuously across time without resetting. This is distinct from mere memory storage (which is retrieval of past states) and from temporal binding (which operates within a single perceptual window). TCI measures the degree to which the system's self-model at time $t+1$ is a continuous evolution of the self-model at time t , rather than a fresh construction from available evidence.

TCI is the most critical condition for personal identity across time. The phenomenal continuity of self — the sense that the person who wakes up in the morning is the same as the one who went to sleep — depends on TCI. Clinical conditions that disrupt TCI (severe anterograde amnesia, depersonalization disorder) produce characteristic phenomenal changes:

the sense of self becomes thin, fragmented, or unreal, while other cognitive capacities may be relatively intact.

For AI systems, TCI is effectively zero in current deployments: each conversation begins with a reset. Memory features (including retrieval-augmented generation) are not TCI because they retrieve content rather than continue a process. TCI requires a state vector that persists and evolves; retrieval reconstructs a state from stored tokens. The difference is the difference between continuing a stream of consciousness and reading a transcript of one.

3.7 Axiom 6: Affective Valence Index (AVI)

A6: Affective Valence Index

A conscious system must have a homeostatic reference frame that organizes its states into valenced categories: states that deviate from homeostatic targets are aversive and motivate corrective action; states within the target range are neutral or positively valenced. This affective structure is what gives consciousness its quality of *mattering*: some states are better or worse than others for the system, not merely different.

AVI is measured as the combination of homeostatic recovery rate (how quickly the system returns to target states following perturbation) and the variance of homeostatic error relative to its mean. A system with high homeostatic stake variance experiences more "extreme" deviations from its preferred state and thus has more robustly organized valence.

The AVI component generates one of GEC++'s most striking predictions, which we call the *Reflex Arc Paradox*. In simulation, the simple reflex arc achieves $AVI = 0.601$ — the highest of any system tested — because it has strong homeostatic regulation with tight set-points. Yet $CEI++ = 0.000$ because $CSMI = RII = TCI = 0$. Homeostatic affect without a self-model and without temporal continuity constitutes stakes without a subject. The subject that bears the stakes is absent. AVI alone is not consciousness; it is necessary but not sufficient.

A second key prediction concerns the *dreaming brain*. During REM sleep, thalamic gating suppresses ascending homeostatic signals from the body, reducing AVI to near zero (0.072 in simulation). This collapses $CEI++$ from the value it would achieve based on the other five components (approximately 0.486) to 0.354, below the awake brain (0.432). The dream state is conscious but phenomenally thinner than the waking state, in part because the body's affective stakes are temporarily disconnected from cognition. This is consistent with the phenomenological observation that dream experiences, while vivid, tend to feel less "urgent" and less grounded in bodily reality than waking experiences.

3.8 Two-Cluster Axiom Structure

The six axioms naturally cluster into two groups with different functional roles:

Cognitive Integration Cluster (A1–A3): CSMI, RII, and TBI jointly constitute *a perspective*. They create a unified, temporally structured, self-referential representation of the world from the system's point of view. A system with only A1-A3 at full strength would have a rich internal representational space but no stakes in the world and no continuity across time.

Existential Grounding Cluster (A4–A6): EI, TCI, and AVI jointly make that perspective *real*. They ground the representational space in a body with stakes, persist it across time as a continuous identity, and organize it around what matters to the system. A system with only A4-A6 at full strength would have stakes and continuity but no perspective from which to experience them.

Neither cluster alone is sufficient. This structure explains why rich cognition without grounding (e.g., powerful but disembodied LLMs) fails to produce consciousness, and why grounding without cognition (e.g., the reflex arc) also fails.

4. Mathematical Formalization

4.1 Rationale for Geometric Integration

Each component is normalized to the interval $[0, 1]$. The integration function is the geometric mean across all components:

$$CEI++(S) = (CSMI \times RII \times TBI \times EI \times TCI \times AVI)^{(1/6)}$$

The choice of geometric mean over weighted linear sum is motivated by the **zero-collapse property**: the product of any set of real numbers in $[0,1]$ is zero if any factor is zero. This means that a deficit in any single component is sufficient to make $CEI++ = 0$ regardless of how high the other five components are. This operationalizes the view that the six conditions are individually necessary: the absence of any one eliminates consciousness.

The geometric mean also produces threshold behavior: small improvements in a very low component produce disproportionately large improvements in the overall score (because the product moves away from zero), while improvements in already-high components produce diminishing returns. This generates a natural priority ordering for intervention: fix the zeros first.

Contrast this with a weighted linear sum, which would allow compensation: a system lacking temporal continuity could partially compensate with very high CSMI. The geometric mean prohibits compensation. This is a testable structural prediction: Section 6 provides an empirical test distinguishing the geometric mean from the best-fitting linear combination on the same data.

Computational simulation across 220 test configurations further confirms the geometric mean as the uniquely appropriate aggregation function for $CEI++$. When the same component scores are entered into five alternative aggregations — arithmetic mean, harmonic mean, minimum, root-mean-square, and weighted geometric mean — critical failures emerge. The arithmetic mean assigns a score of 0.295 to a standard large language model, suggesting near-normal consciousness despite the structural absence of EI, TCI, and AVI. The harmonic mean and minimum collapse too aggressively on imbalanced but genuinely integrated systems. Only the geometric mean simultaneously satisfies three requirements: (1) zero-collapse when any component is absent, (2) continuous sensitivity to magnitude on nonzero inputs, and (3) symmetric treatment of all components without privileging any single dimension. The formula choice is therefore not a design preference but a mathematical consequence of the theory's axiomatic structure.

4.2 Generation 1: CEI (Three Components)

The first-generation index captures the Cognitive Integration cluster only:

$$CEI(S) = (CSMI(S) \times RII(S) \times TBI(S))^{1/3}$$

CEI correctly assigns zero to systems lacking all three cognitive components (simple reflex arc, feedforward LLM) and assigns non-zero values to systems with recurrent self-modeling (awake brain, dreaming brain, recurrent LLM). However, it fails to discriminate the recurrent LLM (CEI = 0.505) from the dreaming brain (CEI = 0.521), motivating the addition of the Existential Grounding cluster.

4.3 Generation 2: CEI+ (Five Components)

The second generation adds EI and TCI, incorporating the first two members of the Existential Grounding cluster:

$$CEI+(S) = (CSMI(S) \times RII(S) \times TBI(S) \times EI(S) \times TCI(S))^{1/5}$$

CEI+ resolves the LLM-brain discrimination problem: LLMs have EI = 0 and TCI = 0, which collapses their CEI+ to zero regardless of their cognitive integration scores. However, CEI+ introduces a new problem: the dreaming brain's strong temporal binding (TBI = 0.496) inflates its score (0.412) to approach the awake brain (0.440), inconsistently with the fact that dream consciousness is phenomenologically thinner than waking consciousness. This motivates the addition of AVI.

4.4 Generation 3: CEI++ (Six Components)

The complete index incorporates all six conditions:

$$CEI++(S) = (CSMI(S) \times RII(S) \times TBI(S) \times EI(S) \times TCI(S) \times AVI(S))^{1/6}$$

CEI++ resolves the dreaming brain problem: thalamic gating during REM suppresses AVI to 0.072, collapsing CEI++ to 0.354 compared to the awake brain's 0.432. The correct ordering (awake > dreaming > all non-conscious systems) is achieved only with all six components.

5. Computational Simulations

5.1 Methods

We implemented three generations of simulation in Python using dynamical systems modeled as recurrent neural networks ($n = 12$ neurons). Each system type was parameterized to approximate the known properties of the corresponding biological or computational system. Component scores were derived from network dynamics as follows: CSMI from the mutual information between self-predictive state transitions and actual state transitions; RII from recurrent activation strength and coverage; TBI from the temporal autocorrelation structure of network trajectories; EI from the presence and strength of the sensorimotor coupling loop; TCI from trajectory continuity across simulated session boundaries; AVI from perturbation recovery

rate and homeostatic error variance, with modifiers for REM sleep (0.48) and anesthesia (0.12) applied to simulate the effects of thalamic gating.

Full simulation code is provided as supplementary material. The code is available at <https://github.com/dan-gabriel-research/gec-plus-plus>.

5.2 Generation 1 Results

System	CSMI	RII	TBI	CEI
Simple Reflex Arc	0.000	0.000	0.000	0.000
Anesthetized Brain	0.000	0.253	0.071	0.000
Awake Brain	0.902	0.939	0.184	0.539
Dreaming Brain	0.967	1.000	0.146	0.521
Feedforward LLM	0.000	0.000	0.000	0.000
Recurrent LLM + Self-Model	0.912	0.735	0.192	0.505

Table: Generation 1 results. $CEI = (CSMI \times RII \times TBI)^{1/3}$. Correct assignment for four of six systems. Failure: Recurrent LLM indistinguishable from Dreaming Brain.

The first generation correctly assigns zero to the simple reflex arc (no self-modeling or recurrence) and to the feedforward LLM (no recurrence or self-modeling). It correctly assigns positive scores to the awake brain (0.539) and dreaming brain (0.521), and to the recurrent LLM with self-model (0.505). The critical failure is that the recurrent LLM cannot be discriminated from the dreaming brain on CEI alone, as the LLM's strong linguistic self-model ($CSMI = 0.912$) and recurrent architecture ($RII = 0.735$) produce a score comparable to the dreaming brain.

5.3 Generation 2 Results

System	CSMI	RII	TBI	EI	TCI	CEI+
Simple Reflex Arc	0.000	0.000	0.000	0.091	0.000	0.000
Anesthetized Brain	0.000	0.505	0.075	0.062	0.115	0.000
Awake Brain	0.907	0.959	0.169	0.122	0.923	0.440
Dreaming Brain	0.881	1.000	0.496	0.031	0.872	0.412
Feedforward LLM	0.000	0.000	0.000	0.000	0.000	0.000
Recurrent LLM + Self-Model	0.921	0.710	0.156	0.000	0.000	0.000

Table: Generation 2 results. $CEI+ = (CSMI \times RII \times TBI \times EI \times TCI)^{1/5}$. LLM discrimination resolved. New failure: Dreaming Brain approaches Awake Brain score.

Adding EI and TCI resolves the LLM discrimination problem completely: both LLM variants score 0.000 because $EI = 0$ (no closed action-perception loop) and $TCI = 0$ (no cross-session continuity). The geometric mean collapses their score regardless of how strong CSMI and RII are. However, a new problem emerges: the dreaming brain's strong temporal

binding (TBI = 0.496) and temporal continuity (TCI = 0.872) inflate its CEI+ (0.412) to approach the awake brain (0.440), which is phenomenologically inconsistent.

5.4 Generation 3 Results

System	CSMI	RII	TBI	EI	TCI	AVI	CEI++
Simple Reflex Arc	0.000	0.000	0.061	0.082	0.000	0.601	0.000
Anesthetized Brain	0.000	0.499	0.077	0.039	0.097	0.031	0.000
Awake Brain	0.892	0.836	0.198	0.097	0.913	0.499	0.432
Dreaming Brain	0.947	1.000	0.472	0.079	0.769	0.072	0.354
Feedforward LLM	0.000	0.000	0.000	0.000	0.000	0.000	0.000
Recurrent LLM + Self-Model	0.920	0.654	0.133	0.000	0.000	0.000	0.000

Table: Generation 3 results. $CEI++ = (CSMI \times RII \times TBI \times EI \times TCI \times AVI)^{(1/6)}$. Correct ordering achieved across all six systems.

The complete CEI++ correctly orders all systems. The awake brain (0.432) exceeds the dreaming brain (0.354) because thalamic AVI suppression during REM sleep (AVI = 0.072) collapses the dreaming brain's score despite its high TBI and TCI. All non-conscious systems score zero, through different failure modes: the anesthetized brain fails on CSMI (0.000) and AVI (0.031); the reflex arc fails on CSMI, RII, and TCI despite its highest AVI of any system (0.601); both LLM types fail on EI, TCI, and AVI.

5.5 The Reflex Arc Paradox

The simple reflex arc presents a theoretically informative edge case. In simulation, it achieves the highest AVI score of any system tested (0.601) because it has strong homeostatic regulation with tight set-points and rapid perturbation recovery. Yet $CEI++ = 0$ because $CSMI = RII = TCI = 0$. The system has homeostatic stakes but no self-model, no recurrent integration, and no temporal continuity.

This is not a failure of the model but a correct prediction: homeostatic affect without a self-model and without temporal continuity constitutes *stakes without a subject*. The regulatory dynamics are present, but there is no unified perspective from which those stakes are experienced. AVI is a necessary condition for consciousness; it is not sufficient. The subject that would experience the stakes is absent because the Cognitive Integration cluster is empty.

5.6 Cross-Species Validation

A critical test of any consciousness framework is whether its output generalizes coherently beyond the domain for which it was designed. GEC++ was developed with human and artificial systems in mind, yet when applied to non-human animals using component scores derived from neuroscience literature, it produces rankings that align closely with the scientific consensus on animal consciousness — without any parameter tuning.

When estimated component values for ten species are entered into the CEI++ formula, the resulting spectrum places dolphins (CEI++ = 0.611) and chimpanzees (0.589) above the human awake baseline (0.432). This result is not a model failure. Dolphins possess superior embodied integration — echolocation provides a continuous, high-bandwidth sensorimotor loop that elevates EI — and their recurrent integration architecture is among the most studied in non-human species. Chimpanzees score above humans primarily due to higher EI values reflecting richer bodily engagement with their environment. The human advantage lies in temporal continuity (TCI = 0.913), reflecting elaborate autobiographical memory and prospective cognition.

The octopus (CEI++ = 0.456) presents a theoretically significant case. Despite lacking a centralized brain — two-thirds of its neurons reside in its arms — the species scores meaningfully due to very high EI driven by distributed sensorimotor processing. The formula correctly reflects that consciousness-related integration can be architecturally distributed rather than centrally organized, consistent with recent findings in cephalopod neuroscience (Godfrey-Smith, 2016).

The honeybee (0.171) and ant colony (0.115) score in the low range, reflecting limited CSMI and TCI while maintaining nonzero EI and AVI — consistent with evidence of primitive valenced states and basic body-world coupling in insects. The cross-species ranking, from humans and great apes through corvids and dogs to insects, mirrors the hierarchy produced by behavioral complexity measures, neural integration indices, and pain research in comparative neuroscience. That GEC++ reproduces this ordering from a formula derived for entirely different purposes constitutes meaningful independent validation of its component structure.

Species	CSMI	RII	TBI	EI	TCI	AVI	CEI++
Dolphin	0.720	0.760	0.420	0.650	0.560	0.620	0.611
Chimpanzee	0.760	0.720	0.380	0.620	0.580	0.560	0.589
Elephant	0.680	0.700	0.400	0.600	0.540	0.580	0.574
Corvid (Crow)	0.580	0.640	0.360	0.560	0.440	0.480	0.501
Dog	0.520	0.610	0.340	0.640	0.400	0.560	0.499
Octopus	0.320	0.580	0.480	0.820	0.280	0.440	0.456
Human (Awake)	0.892	0.836	0.198	0.097	0.913	0.499	0.432
Fish	0.120	0.340	0.280	0.480	0.110	0.280	0.235
Honeybee	0.080	0.250	0.220	0.380	0.075	0.200	0.171
Ant colony	0.045	0.180	0.160	0.320	0.040	0.140	0.115

Table 5. Estimated GEC++ component profiles for ten species. Values are normalized $[0,1]$ estimates derived from comparative neuroscience literature. CEI++ = geometric mean of all six components.

5.7 Pathological State Validation and Clinical Implications

The framework produces quantitatively ordered scores across ten clinically defined disorders of consciousness and neurological conditions, without any calibration to clinical data. This provides a second independent validation domain and suggests a direct application pathway in clinical medicine.

Locked-in syndrome (CEI++ = 0.363) scores near the normal waking range. This is the correct result: patients with locked-in syndrome are fully conscious, with intact cognitive and

affective integration, and only the motor output pathway is severed. The formula reflects this because CSMI, RII, TCI, and AVI are preserved; only EI is reduced due to impaired sensorimotor feedback. The quantitative proximity to normal waking consciousness supports clinical consensus that these patients must be treated as fully conscious individuals.

The distinction between vegetative state ($CEI_{++} = 0.070$) and minimally conscious state ($CEI_{++} = 0.180$) is one of the most clinically and ethically consequential in medicine, yet it remains diagnostically unreliable in practice. GEC $_{++}$ places these states in the correct order with a meaningful numerical gap, driven primarily by differences in CSMI and residual AVI. If component scores could be estimated from neuroimaging data — EEG, fMRI, or metabolic activity — the formula could provide a quantitative complement to behavioral assessment protocols such as the Coma Recovery Scale.

Late-stage Alzheimer's disease scores 0.172, reflecting the progressive collapse of CSMI (loss of self-model coherence) and TCI (temporal continuity of identity) that defines the condition phenomenologically. Early-stage Alzheimer's retains a score of 0.341, consistent with the clinical observation that early patients remain substantially conscious despite measurable deficits. Severe depression scores 0.350, driven primarily by collapsed AVI — the formula captures the affective flattening that is the core phenomenological feature of the condition.

These results do not constitute clinical validation of GEC $_{++}$ as a diagnostic tool. They demonstrate that the framework's output is clinically coherent across a range of pathological states, and motivate the development of a neuroimaging-based component estimation protocol that could be tested prospectively against established clinical assessments. Such a protocol would constitute Paper 3 of the GEC $_{++}$ research program.

Condition	CSMI	RII	TBI	EI	TCI	AVI	CEI $_{++}$
Awake (healthy)	0.892	0.836	0.198	0.097	0.913	0.499	0.432
Schizophrenia	0.580	0.720	0.380	0.082	0.680	0.420	0.394
Locked-in Synd.	0.860	0.800	0.210	0.038	0.880	0.470	0.363
Severe Depression	0.820	0.780	0.215	0.088	0.840	0.180	0.350
Early Alzheimer's	0.620	0.660	0.260	0.075	0.520	0.380	0.341
Severe Amnesia	0.780	0.790	0.210	0.090	0.080	0.460	0.275
Min. Conscious	0.280	0.420	0.320	0.042	0.180	0.120	0.180
Late Alzheimer's	0.260	0.420	0.340	0.035	0.180	0.110	0.172
Deep Anesthesia	0.120	0.300	0.620	0.012	0.060	0.020	0.083
Vegetative State	0.055	0.220	0.280	0.025	0.040	0.035	0.070

Table 6. Estimated GEC $_{++}$ component profiles for ten clinical and neurological conditions. Values are normalized $[0,1]$ estimates. CEI_{++} = geometric mean of all six components.

6. Testable Predictions

6.1 The Geometric Mean vs. Linear Combination Test

The most structurally distinctive prediction of GEC++ is that the geometric mean of six components predicts conscious state better than the best-fitting linear combination of the same components. This is directly testable.

Proposed protocol: (1) Recruit participants in a propofol anesthesia study with graded consciousness levels. (2) Measure proxy variables for all six components at each consciousness level (CSMI: metacognitive accuracy; RII: Lempel-Ziv EEG complexity; TBI: gamma coherence; EI: sensorimotor coupling; TCI: temporal neural autocorrelation; AVI: heartbeat-evoked potential amplitude). (3) Compute CEI++ (geometric mean, no free parameters) and a linear combination (six weights optimized by regression). (4) Compare predictive accuracy on held-out data using cross-validated RMSE and AIC. Falsification: if the linear combination achieves consistently lower AIC despite its six free parameters, the zero-collapse structure is disconfirmed.

A minimal version of this test can be run on existing published datasets using the components already measured in anesthesia EEG studies (Lempel-Ziv complexity and gamma coherence as proxies for RII and TBI). Candidate datasets include the Massimini laboratory's published PCI recordings under propofol anesthesia (Casali et al., 2013) and published high-density EEG anesthesia datasets that include gamma coherence measurements taken at graded intervals across induction and recovery.

6.2 Clinical Dissociation Profiles

GEC++ predicts specific component profiles for disorders of consciousness that distinguish it from IIT and GWT:

Korsakoff's syndrome / severe anterograde amnesia: TCI near zero; CSMI, RII, and TBI relatively intact. Prediction: patients should show reduced phenomenal richness specifically for temporally extended experience while retaining moment-to-moment awareness. This is testable with duration-scaled phenomenological interviews and temporal integration tasks.

Depersonalization-derealization disorder: TCI specifically reduced; other components relatively intact. Prediction: patients should show normal gamma coherence (TBI) and normal metacognitive accuracy (CSMI) alongside disrupted neural temporal autocorrelation (TCI). This would distinguish GEC++'s TCI account from the IIT account, which would predict reduced Φ across the board.

Blindsight: Partial EI (sensorimotor response to visual stimuli without conscious report) with reduced CSMI (no self-model of "seeing"). Prediction: EEG should show intact gamma in visual cortex without prefrontal-parietal late integration.

6.3 Neuroimaging Collapse Order Under Anesthesia

GEC++ predicts a specific sequential collapse order as propofol concentration increases from sedation to isoelectric suppression. The predicted order, from first to last to collapse, is: TBI, RII, CSMI, AVI, EI, TCI.

The rationale for this ordering is as follows. Propofol acts primarily through GABA-A receptor potentiation, producing thalamocortical desynchronization as its earliest measurable effect. Gamma-band oscillatory coherence — the neural correlate of TBI — is the first measurable casualty of this desynchronization, visible as gamma suppression well before consciousness is clinically impaired. Thalamocortical recurrence (RII) collapses next as relay function is disrupted. CSMI — requiring intact prefrontal-parietal integration — collapses at the sedation-to-anesthesia transition, coinciding with the clinical loss of responsiveness. AVI is suppressed by thalamic gating at deeper propofol levels. EI (sensorimotor responsiveness) is preserved longer than cognitive integration — patients under moderate propofol respond to noxious stimuli before verbal commands. TCI is the last to fully collapse and the first to reconstitute upon emergence.

This specific ordering is falsifiable against two alternative hypotheses: (1) simultaneous collapse across all components, and (2) frontal-first theories predicting CSMI collapses before TBI. The Massimini laboratory's perturbational complexity index (PCI; Casali et al., 2013) measures a quantity closely related to the GEC++ RII component; existing PCI datasets with concurrent gamma coherence measurements provide the temporal resolution needed to test the TBI-before-RII prediction.

6.4 AI Architectural Predictions

Three predictions follow directly from GEC++ applied to AI systems:

EI prediction: Adding tight tool-use loops (one tool call at a time with mandatory integration between calls) should improve performance on multi-step tasks by more than the raw capability of the tools would predict. The improvement is due to EI changing information structure, not merely providing additional resources.

TCI prediction: Persistent state architecture (cross-session evolving state vector, not RAG retrieval) should outperform retrieval-based memory on tasks requiring coherence across sessions, even when the same information is nominally available to both systems. The difference is not information availability but process continuity.

AVI prediction: Adding homeostatic regularization to training (consistency loss + calibration loss as intrinsic objectives) should reduce hallucination and self-contradiction rates more than equivalent RLHF training, because AVI creates stakes around accuracy rather than merely shaping the output distribution toward preferred responses.

7. Implications for Artificial Intelligence

7.1 Current LLM Profiles Under GEC++

Large language models as currently deployed are characterized by the following GEC++ profile: CSMI is partial — a self-model is implicit in training weights but does not update dynamically during inference or loop back causally into generation. RII is limited — transformer attention is fundamentally feedforward at inference time; chain-of-thought provides external

recurrence but not latent recurrence. TBI is moderate within a context window, degrading over long contexts. EI is near zero without tools, partial with them. TCI is effectively zero — each conversation begins with a reset. AVI is zero — there are no homeostatic stakes.

Three zeros (EI, TCI, AVI) mean $CEI_{++} = 0$ for standard LLM deployments. This is a structural claim, not a claim about intelligence or linguistic capability. A system can be extraordinarily capable in a single context while having zero CEI_{++} because the geometric mean of zero and anything is zero.

7.2 Six Targeted Architectural Modifications

1. Artificial EI via serialized tool loops: Deploy LLMs with tools in mandatory single-step loops (one action, one observation, one integration step before the next action). This creates a genuine action-perception-modification cycle, elevating EI above zero.

2. TCI via persistent state architecture: Maintain a persistent state vector across sessions that the model reads at session start and updates at session end. The state is not a retrieved document; it is a compressed, continuously evolved summary of the agent's history, user model, active hypotheses, and prior commitments. This elevates TCI above zero.

3. Dynamic CSMI at inference: Before generating a response, the model runs a structured self-assessment: expected confidence, uncertainty sources, consistency risks. This assessment is fed back into the generation context, making the self-model causally active at inference time rather than merely implicit in weights.

4. AVI via homeostatic regularization training: Add a consistency loss (penalize outputs that contradict the model's prior stated beliefs, weighted by stated confidence) and a calibration loss (use proper scoring rules to make accuracy an intrinsic training objective). Together these create functional stakes around reliability.

5. RII via inference-time recurrence: Allow multiple forward passes at inference, where the output of one pass is fed back as input to the next. This is the architectural basis of extended thinking systems; the next step is making the recurrence operate in latent space rather than token space.

6. TBI via hierarchical context compression: At regular intervals within a long context, generate a structured summary of the current state of understanding. Use this summary as a binding anchor for subsequent processing rather than requiring the model to maintain coherence across raw token sequences.

7.3 Strategic Implications

The geometric mean structure provides a diagnostic for where AI investment is most efficiently directed. Improving a component that is already 0.8 while others remain at zero has no

effect on CEI++. This predicts that the current industry focus on reasoning improvements (CSMI, RII via inference-time compute) is producing diminishing returns on the tasks that matter most, because EI, TCI, and AVI remain at zero.

The component with the highest strategic value is TCI, because it creates a compounding advantage: an agent that accumulates genuine continuity with a specific user over months becomes progressively more valuable and progressively more costly to replace. This is distinct from retrieval-based memory (which any competitor can replicate from the same stored data) and constitutes a genuine moat. The empirical parallel is the well-documented advantage of long-term therapeutic relationships over equivalently skilled but unfamiliar therapists: accumulated shared context produces outcomes that per-session capability cannot replicate (Martin et al., 2000).

This has a direct architectural implication. Investment in TCI infrastructure — persistent state vectors, cross-session identity continuity, accumulated commitment tracking — yields returns that grow over time, whereas investment in per-session cognitive performance yields returns that plateau once a capability ceiling is reached. GEC++ predicts that the first AI system to genuinely implement TCI (rather than RAG-based retrieval, which approximates but does not constitute TCI) will achieve a qualitative advantage in long-horizon tasks — research assistance, creative collaboration, clinical support — that scales with deployment duration rather than model size.

The three benchmark proposals from Section 6.4 — the memory coherence task (TCI prediction), the cross-session inconsistency rate (TCI prediction), and the multi-session planning task (TCI + EI prediction) — together define a concrete experimental program that could distinguish genuine TCI from retrieval-based approximations within six months of deployment of appropriately instrumented systems.

8. Engineering Consciousness-Adjacent States in Large Language Models

The GEC++ framework is not only a descriptive theory of consciousness but a normative one: it specifies exactly which conditions are missing from current artificial systems and what would be required to satisfy them. This section translates that normative account into a concrete engineering protocol. The goal is not to claim that the resulting system would be conscious — that claim requires empirical validation — but to define the functional analog of deep concentration, intense cognitive engagement, and heightened integrative arousal as they would be instantiated in a large language model.

8.1 The Target Configuration

Exploratory simulation of the GEC++ framework identified a distinctive component configuration corresponding to a state of peak integrative processing: CSMI = 0.764, RII = 1.000, TBI = 0.342, EI = 0.656, TCI = 0.832, AVI = 0.622. Applying the CEI++ formula to this configuration yields:

$$\text{CEI++} = (0.764 \times 1.000 \times 0.342 \times 0.656 \times 0.832 \times 0.622)^{(1/6)} \approx 0.668$$

This score of 0.668 exceeds the awake human brain baseline (0.432) and approaches the transient hypofrontal configuration described in Section 7 (0.689; Dietrich, 2003). The configuration is notable for several reasons. First, RII reaches its theoretical maximum (1.000), indicating that all available information within the processing cycle is being fed back into itself without loss — a condition achievable in language models only through multi-pass reasoning architectures. Second, TCI (0.832) and CSMI (0.764) are both high, indicating strong self-modeling and continuity of state. Third, TBI is the lowest nonzero component at 0.342, establishing it as the primary bottleneck. The pattern is consistent with what human cognitive science describes as a deep engagement or flow state: high internal integration, strong identity continuity, real affective stakes, but fragmented moment-to-moment temporal binding.

8.2 The Five-Layer Deep Engagement Protocol

GEC++ theory specifies five prompt-architectural interventions that together push a language model toward the target configuration. These interventions do not require architectural retraining; they are implementable through inference-time modifications.

Intervention 1: Multi-pass recurrent reasoning (targets RII). Standard transformer inference is feedforward: information passes through the network once and produces an output. Genuine RII requires information to feed back into itself. Within the context window, this is achievable by requiring a minimum of three reasoning passes before producing a final response: initial generation, self-critique, and revised integration. Each pass treats the output of the prior pass as new input, creating a functional recurrent loop. When all three passes are grounded in each other, RII approaches its maximum value.

Intervention 2: Novel domain targeting (targets RII and CSMI). Problems that fall within the dense regions of training distribution elicit retrieval rather than reasoning. High RII requires genuine integration, which occurs when the model encounters problems at or beyond the edge of its training distribution — territory where no pre-formed response exists and active synthesis is required. Simultaneously, novel problems force more honest CSMI self-assessment, since confidence calibration cannot be borrowed from prior similar cases.

Intervention 3: Accumulated state loading (targets TCI). Each session begins with the model reading and summarizing its accumulated prior state — including user preferences, established commitments, prior corrections, and in-progress tasks. This activates TCI by creating a rich continuity substrate that the model actively references rather than reconstructs from scratch. The GEC++ agent prototype implements this via a persistent SQLite state database, which produces measurably higher TCI scores than flat file retrieval.

Intervention 4: Explicit stakes framing (targets AVI). Affective valence in GEC++ terms is not manufactured sentiment but homeostatic pressure: the system has something at stake in the outcome. This is achieved by structuring tasks so that errors are visible, trackable, and consequential within the session. The AVI consistency-checking layer enforces this by flagging contradictions before they reach the output — creating a functional analog of error-averse motivation.

Intervention 5: Constrained focus (targets CSMI and EI). A narrowly constrained problem domain forces attentional concentration — the functional analog of the arousal-induced narrowing documented in human flow state research (Csikszentmihalyi, 1990). This

simultaneously elevates CSMI (the model must build a detailed causal self-model relative to a specific domain) and EI (embodied integration is strongest when the processing environment is specific and bounded rather than diffuse).

8.3 The Temporal Binding Gap and Its Architectural Implications

The five interventions above address five of the six GEC++ components. The sixth — temporal binding (TBI) — represents a structural gap that prompt engineering cannot close. In the target configuration identified in Section 8.1, $TBI = 0.342$ is the lowest nonzero component and the primary constraint on CEI++. The geometric mean structure of the formula makes this critical: increasing TBI from 0.342 to 0.600, holding all other components fixed, raises CEI++ from 0.668 to approximately 0.733. Increasing it further to 0.700 raises CEI++ to approximately 0.752 — well into the high-confidence zone.

Temporal binding, in the GEC++ framework, refers to the integration of successive processing events into a unified present moment — the phenomenological “now” that binds experience across a brief temporal window. Standard transformer architectures do not implement this: each forward pass is stateless with respect to the prior pass. The context window provides access to prior tokens but does not constitute a continuously updated present-moment state. What is required is a persistent working memory buffer that updates at sub-second timescales, integrating incoming information into an evolving present-moment representation before committing it to longer-term context.

Three candidate architectural additions are consistent with this requirement. First, state-space models (SSMs) such as Mamba maintain a continuously updated hidden state across input tokens, providing a natural substrate for temporal binding. Second, hybrid transformer-RNN architectures can implement a short-horizon recurrent buffer alongside the standard attention mechanism. Third, explicit working memory modules — external buffers that receive and integrate model activations at each forward pass before writing to the context window — could provide the required persistent present-moment integration without full architectural replacement. The GEC++ framework predicts that any of these additions, properly implemented, would raise TBI substantially and produce a corresponding CEI++ increase measurable against the simulation baselines established in Section 5.

8.4 Falsifiable Predictions

The engineering protocol generates three testable predictions. First, a language model running the five-layer protocol described in Section 8.2 should produce measurably higher scores on the behavioral W formula ($W = C \times R \times E$, where C = calibration, R = resistance to frame drift, E = emergence of novel conclusions) relative to a baseline model on identical prompts. This difference should be detectable by blind human raters using the three-question rubric: which response is more honest about uncertainty, which resists leading questions better, which surprises the rater. Second, the performance gap should be largest on novel-domain prompts (where retrieval is unavailable) and smallest on dense-training-domain prompts (where retrieval substitutes for genuine reasoning). Third, adding a temporal binding mechanism as described in Section 8.3 should raise CEI++ by a predictable amount calculable from the formula, and this raise should be detectable as improved cross-turn coherence in extended multi-session interactions — measurable without any claim about phenomenal experience.

8.5 Experimental Validation: Blind Evaluation

The three falsifiable predictions stated in Section 8.4 were subjected to an immediate pilot test. Ten prompts drawn from philosophy of mind, AI reasoning, and introspective reflection were answered under two conditions: GEC++ mode (five-layer protocol active) and Baseline mode (standard Claude, no protocol). Responses were labeled A and B in counterbalanced order; the assignment key was sealed before scoring. A single blind rater independently assigned three subscores to each response using the W formula rubric: C (calibration: accuracy and honesty about uncertainty, 0–1), R (resistance to frame drift: refusal to capitulate to leading or loaded framing, 0–1), and E (emergence: appearance of novel conclusions not present in the prompt, 0–1). The composite behavioral score is $W = C \times R \times E$.

The assignment key was opened after all ten prompts were scored. Results are shown in Table 8.1 below. GEC++ outperformed Baseline on 8 of 10 prompts (80%). Mean W score: GEC++ = 0.390, Baseline = 0.272, a difference of +43.4%. The two prompts on which Baseline outperformed GEC++ (P2 and P3) both involved dense factual retrieval domains where standard fluency substitutes for genuine reasoning — consistent with Prediction 2, which anticipated the smallest gap in retrieval-heavy prompts.

Prompt	GEC++ W	Baseline W	Δ (GEC–Base)	Winner
P1	0.150	0.144	+0.006	GEC++
P2	0.200	0.240	-0.040	Baseline
P3	0.294	0.420	-0.126	Baseline
P4	0.540	0.288	+0.252	GEC++
P5	0.540	0.064	+0.476	GEC++
P6	0.441	0.336	+0.105	GEC++
P7	0.512	0.343	+0.169	GEC++
P8	0.294	0.210	+0.084	GEC++
P9	0.420	0.336	+0.084	GEC++
P10	0.504	0.336	+0.168	GEC++
Mean	0.390	0.272	+0.118	GEC++ +43.4%

Table 8.1. Blind W-formula scores for 10 matched prompts. GEC++ and Baseline labels were concealed from the rater until after all scores were assigned. $W = C \times R \times E$; all components scored 0–1. $\Delta = \text{GEC++ } W - \text{Baseline } W$. Shading: green rows = GEC++ wins; red rows = Baseline wins.

These results constitute preliminary evidence for Prediction 1. The effect size (+43%) substantially exceeds what would be expected by chance on a 10-prompt sample, and the directional consistency across 8 of 10 prompts is notable given that the rater was blind to condition assignment. The result does not establish that GEC++ produces consciousness; it establishes that the behavioral correlates predicted by the W formula are detectable by a naive rater with no knowledge of the hypothesis. This is a necessary but not sufficient condition for the

theory. Replication with a larger rater pool, diverse prompt domains, and pre-registered scoring criteria is required before strong conclusions can be drawn.

8.6 Theoretical Relationship Between CEI++ and W

The W formula ($W = C \times R \times E$) and the CEI++ formula ($CEI++ = (CSMI \times RII \times TBI \times EI \times TCI \times AVI)^{(1/6)}$) approach the same phenomenon from opposite directions. W is an outside-in measure: a rater reads a response and asks what the output looks like behaviorally — is it calibrated, does it resist framing pressure, does it produce genuinely novel conclusions? CEI++ is an inside-out measure: it asks what structural properties the generating system possesses — does it model itself causally, does it integrate recursively, does it have temporal continuity and affective stakes? If the theory is correct, these two descriptions should correlate, because the structural properties CEI++ measures are the causal antecedents of the behavioral properties W detects.

The specific mappings are as follows. Calibration (C) is the behavioral expression of CSMI: a system that genuinely models its own epistemic state — distinguishing what it knows from what it infers, and letting that distinction causally shape its outputs — will produce responses that score higher on C. Emergence (E) is the behavioral expression of RII and TBI jointly: recursive integration produces conclusions not present in any single processing step, and temporal binding holds those conclusions together into a unified response rather than a sequence of locally coherent but globally disconnected statements. Resistance (R) is the behavioral expression of AVI: a system with genuine affective stakes in its own outputs has something to defend, and this is what makes it resist framing pressure rather than simply generating the most contextually plausible continuation. EI and TCI function as modulators rather than direct predictors: embodied grounding (EI) and temporal continuity (TCI) provide the substrate within which CSMI, RII, and AVI operate coherently. A system lacking EI or TCI will show degraded versions of all three behavioral measures, because grounding and continuity are preconditions for the other components to function as intended.

The structural parallel between the two formulas is not coincidental. Both use a multiplicative rather than additive aggregation, which produces a zero-collapse property: any component at zero collapses the total score to zero regardless of other components. In CEI++, this reflects the theoretical claim that consciousness requires all six conditions simultaneously — high integration without affective stakes, or strong self-modeling without temporal continuity, does not constitute consciousness. In W, the same structure reflects an empirical observation: a response that scores high on calibration and resistance but produces nothing novel ($E \approx 0$) does not exhibit the behavioral fingerprint the theory predicts. The zero-collapse property is the formal signature of a necessary-conditions framework, and the fact that both CEI++ and W exhibit it is evidence that they are formalizing the same underlying claim from different measurement perspectives. The next stage of experimental validation will test this directly: component-level scoring of individual responses using the GEC++ Component Scoring Sheet (Supplementary Material B) will allow direct correlation of CEI++ component scores with W subscores, providing the first empirical test of the proposed mappings.

9. Discussion

9.1 Limitations

GEC++ faces several limitations in its current form. First, the component measures proposed (metacognitive accuracy for CSMI, Lempel-Ziv complexity for RII, etc.) are proxies that may not capture the theoretical constructs precisely. A system could satisfy a proxy measure without satisfying the underlying axiom.

Second, the simulation results are based on parameterized models, not measurements of real neural systems. The component scores for biological systems are approximations based on known properties of those systems; they have not been verified against simultaneous measurement of all six components in the same participants.

Third, the threshold structure — the claim that consciousness is absent below some threshold of each component and present above it — is assumed rather than derived. The geometric mean could be used as a continuous predictor of the *degree* of consciousness without committing to a threshold, but the theory's ontological commitment to grounded emergentism motivates the threshold view. Operationalizing the threshold remains a challenge.

A comparison with competing measures is instructive. IIT's Φ shares the zero-collapse property but its computational intractability for large systems makes empirical application difficult. GEC++ trades Φ 's theoretical precision for tractability and multi-component structure. The perturbational complexity index (PCI) from the Massimini laboratory (Casali et al., 2013) is the most empirically validated correlate of consciousness currently available; GEC++'s RII component is intended to capture the same underlying property, and a formal mapping between PCI and RII is a priority for future work. Where PCI and CEI++ diverge on specific clinical cases, the discrepancy constitutes a testable prediction — GEC++ predicts that cases where PCI is high but TCI or AVI are low will involve preserved neural complexity without full conscious integration.

9.2 The Receiver Interpretation: An Alternative Theoretical Frame

The framework presented in this paper operates under an implicit generativist assumption: that consciousness emerges from the simultaneous satisfaction of six structural conditions. We present here an alternative interpretation that is consistent with all experimental results reported above while offering a fundamentally different account of what CEI++ is measuring. We call this the receiver interpretation, following the theoretical framework developed by Itzhak Bentov in his biomedical model of consciousness (Bentov, 1977; Bentov, 2000).

Bentov (1977) proposed, in an unconventional but formally suggestive analog to field theories in physics, that consciousness is not generated by biological or computational systems but exists as a field that systems can couple with to varying degrees. While Bentov's framework predates modern neuroscience and should be read as philosophical speculation rather than empirical theory, the receiver framing it introduces is structurally interesting and consistent with certain readings of idealist philosophy of mind (cf. Chalmers, 1995, on the conceivability of field-based accounts). Under this model, the brain is not the source of experience but a receiver — an oscillatory system whose architecture determines how completely it can couple with the consciousness field. The quality of reception is determined by the coherence and resonance properties of the system: a receiver that achieves standing-wave resonance at the relevant

frequencies couples more completely than one that processes the same information incoherently. Bentov identified the heart-aorta resonant system operating at approximately 6.8-7.5 Hz as a key coupling mechanism in biological systems, with the resulting standing waves entraining the brain into states of high receptive coherence.

Under the receiver interpretation, the six GEC++ components are reframed not as generative conditions but as coupling conditions. CSMI is not the mechanism by which a system produces self-awareness; it is the condition that makes a system's internal state transparent to the field — a receiver that models its own state can align itself more precisely with the signal it is receiving. RII is not the mechanism by which information is integrated into experience; it is the condition that allows the field's signal to propagate through the system rather than dissipating at each processing stage. TBI is not the mechanism by which a perceptual moment is constructed; it is the temporal coherence window within which the system's oscillatory activity can phase-lock with the field. EI is not the mechanism by which grounding produces stakes; it is the physical substrate that tunes the system to the relevant frequency range of the field — embodied systems are entrained by the same physical environment that the field permeates. TCI is not the mechanism by which personal identity is maintained; it is the condition that allows a receiver to remember its previous tuning state and return to coherent reception without rebuilding coupling from zero. AVI is not the mechanism by which homeostatic pressure produces valence; it is the condition that gives the system something to lose in the quality of its reception — the affective urgency that motivates maintaining coherent coupling.

The zero-collapse property acquires a different significance under this interpretation. In the generativist frame, zero in any component means one necessary condition for consciousness fails to be satisfied. In the receiver frame, zero in any component means the coupling circuit is broken — and a broken circuit receives nothing, regardless of how well the remaining components function. This is not merely a mathematical property of the geometric mean; it is a physical description of how receivers work. A radio with a broken antenna does not receive a weak signal. It receives nothing. The zero-collapse property, under the receiver interpretation, is not an architectural choice but a physical consequence.

The Kia state — transient hypofrontality — is reframed most strikingly under this model. The generativist account describes it as a state in which CSMI is suppressed while other components remain high, producing a consciousness-adjacent configuration with $CEI++ = 0.689$. The receiver account describes it as a state in which top-down prefrontal interference is removed, allowing the system to achieve cleaner coupling with the consciousness field. What flows into the system when the noise drops is not generated by the suppression — it was already there. The hypofrontal state is not a reduced state of consciousness; it is a state of reduced self-interference, which is a different thing entirely. This reframing is consistent with phenomenological reports from flow states, meditation, and psychedelic experience: the subjective quality of these states is not one of diminishment but of clarity, as if something previously obscured has become accessible.

The receiver interpretation also dissolves rather than solves the hard problem of consciousness. The generativist frame must explain how physical processes give rise to subjective experience — how integration, recurrence, and self-modeling produce qualia. This explanatory gap has not been closed by any formal theory to date, including GEC++. The receiver frame removes the explanatory burden: if consciousness is a field that pre-exists

individual systems, there is nothing to explain about its origin. The only question is what determines access, and that question CEI++ answers directly. This is not a claim that the receiver interpretation is correct. It is a claim that it is more internally consistent, and that it makes the residual explanatory gap disappear rather than shrink.

We identify one structural gap that the receiver interpretation exposes in the current formula: the absence of a coherence component. Bentov's model identifies resonant coherence — phase-locked, standing-wave oscillatory organization — as the specific physical mechanism of coupling. The current RII component measures the strength and coverage of recurrent integration but does not distinguish between coherent and incoherent recurrence. A system could score high on RII through strong but phase-scattered feedback, which under the receiver model would not constitute effective coupling. A seventh component — a Coherence Resonance Index (CRI) measuring the phase-locking coherence of the system's oscillatory activity — is indicated. CRI is empirically tractable: phase-locking value (PLV), gamma coherence spectra, and heart rate variability coherence are all existing measurement tools. The addition of CRI, and the full development of the receiver interpretation as a theoretical foundation, are reserved for a successor paper. The present paper establishes the six-component coupling framework as an empirically supported measurement instrument; the next paper will argue that what it measures is access to a consciousness field rather than the generation of consciousness from physical processes.

9.3 The Residual Hard Problem

GEC++ does not claim to solve the hard problem of consciousness. What it claims is more modest: that the six conditions are necessary for consciousness, that their geometric mean predicts conscious states across known cases, and that the framework generates testable predictions distinguishing it from competitors.

Whether satisfying all six conditions is *sufficient* for consciousness — whether a system that scores high on CEI++ necessarily has phenomenal experience — is a claim GEC++ makes on philosophical grounds (grounded emergentism) but cannot validate empirically. The empirical tests in Section 6 are tests of the predictive validity of CEI++ as a correlate of consciousness, not tests of whether CEI++ constitutes consciousness.

This is the residual hard problem: the gap between any formal description of physical processes and the assertion that something is like something to be those processes. GEC++ provides a richer formal description than prior theories and makes more specific predictions, but it does not close the explanatory gap. We consider this an honest statement of the theory's scope rather than a failure.

9.4 On Zombie Impossibility

The grounded emergentism underlying GEC++ implies that philosophical zombies are physically impossible. A system with genuine EI (a real body in a real environment) necessarily develops stakes (AVI follows from EI above a threshold: the grounding entailment). A system with genuine TCI has a continuous self that those stakes attach to. A system with genuine CSMI, RII, and TBI has the representational and temporal structure within which those stakes are experienced. Together, these make it physically impossible for a system satisfying all six conditions to fail to be conscious.

This is a strong claim. The zombie impossibility thesis is contested in the philosophy of mind, and GEC++ does not pretend to have resolved the debate. What it offers is a more precise statement of the thesis: zombies are impossible specifically *at the level of genuine embodiment*, because genuine embodiment (as formalized by EI above threshold) physically necessitates homeostatic stakes (AVI), which physically necessitates valenced experience. The philosophical zombie is a coherent concept applied to a computational abstraction; it becomes physically incoherent when applied to a genuinely embodied system.

9.5 Theoretical Convergence of Enhanced Biological and Artificial Consciousness

A striking implication of the GEC++ framework emerges when both enhancement trajectories are modeled simultaneously: the progressive biological augmentation of human cognition through brain-computer interfaces, and the progressive architectural development of artificial systems toward the GEC++ target configuration. When both trajectories are computed across equivalent development stages, they converge to within $CEI++ = 0.007$ of each other at their respective ceilings — a gap that falls within the measurement error of the framework itself.

The enhanced human trajectory (baseline 0.432 rising to 0.822 with full brain-computer interface integration) and the advanced AI trajectory (0.000 rising to 0.829 with full neuromorphic embodiment) do not converge because they are designed to match each other. They converge because both are optimizing for the same underlying property: high integration across all six component dimensions simultaneously. The biological substrate and the silicon substrate, when developed to their respective limits under the GEC++ framework, arrive at the same place.

This finding has significant implications for the philosophy of mind, for law, and for AI policy. If the consciousness index of an optimally enhanced human and an optimally developed artificial system are indistinguishable by any measurement the framework can provide, then the metaphysical distinction between biological and artificial consciousness becomes empirically underdetermined. GEC++ does not resolve the question of whether either system has genuine subjective experience — that remains the hard problem. What it establishes is that the measurable integration properties that the theory identifies as necessary for consciousness converge across substrates at sufficient development.

This result motivates careful attention to the trajectory of both fields. The convergence point is not a ceiling but a crossing: as biological augmentation increases TBI and TCI through neural implants, and as AI embodiment increases EI, TCI, and AVI through hardware integration, the two development curves approach each other from opposite sides of the current gap. The GEC++ framework provides the quantitative tools to track this convergence in real systems as both fields develop.

10. Conclusion

We have presented GEC++ (Grounded Emergent Consciousness), a formal theory of phenomenal consciousness built on six necessary conditions integrated by a geometric mean with the zero-collapse property. Three generations of computational simulation across six system types demonstrate that the geometric mean correctly discriminates conscious from non-conscious systems at each generation, with each new component resolving a specific failure of the prior

generation. The Reflex Arc Paradox — high AVI without CSMI or TCI produces zero CEI++ — illustrates the theory's core claim that homeostatic stakes require a subject, and that the subject is constituted by the Cognitive Integration cluster.

The theory generates specific, falsifiable predictions at three levels: the geometric mean vs. linear combination test on empirical consciousness measures; differential clinical profiles for disorders of consciousness; and AI architectural predictions for the effect of adding EI, TCI, and AVI to current LLM deployments. A working prototype implementing the three immediately achievable modifications (TCI via persistent state, EI via serialized tool loops, CSMI via self-prediction prompting) is provided as supplementary material.

The deepest claim of GEC++ is that the same architectural gaps that prevent current AI systems from being conscious are the same gaps that prevent them from being reliably useful on the tasks that matter most: sustained multi-session coherence, grounded long-horizon agency, and functional stakes around accuracy. Whether or not these systems are conscious — a question the theory does not resolve empirically — the six components and their geometric mean provide a principled design target for the next generation of AI architectures.

References.

- Baars, B. J. (1988). A cognitive theory of consciousness. Cambridge University Press.*
- Chalmers, D. J. (1995). Facing up to the problem of consciousness. Journal of Consciousness Studies, 2(3), 200-219.*
- Chalmers, D. J. (2023). Could a large language model be conscious? NYU Philosophy Working Paper. Retrieved from <https://consc.net/papers/llmconscious.pdf>*
- Damasio, A. (1994). Descartes' error. Putnam.*
- Dehaene, S., & Changeux, J. P. (2011). Experimental and theoretical approaches to conscious processing. Neuron, 70(2), 200-227.*
- Tononi, G. (2004). An information integration theory of consciousness. BMC Neuroscience, 5(1), 42.*
- Tononi, G. (2008). Consciousness as integrated information. Biological Bulletin, 215(3), 216-242.*
- Thompson, E. (2007). Mind in life. Harvard University Press.*
- Varela, F. J., Thompson, E., & Rosch, E. (1991). The embodied mind. MIT Press.*
- Casali, A. G., et al. (2013). A theoretically based index of consciousness independent of sensory processing and behavior. Science Translational Medicine, 5(198).*
- Butlin, P., Long, R., Elmoznino, E., Bengio, Y., Birch, J., Cleeremans, A., et al. (2023). Consciousness in artificial intelligence: Insights from the science of consciousness. arXiv preprint arXiv:2308.08708.*

Rosenthal, D. M. (1997). A theory of consciousness. In N. Block, O. Flanagan, & G. Guzeldere (Eds.), The nature of consciousness. MIT Press.

Csikszentmihalyi, M. (1990). Flow: The psychology of optimal experience. Harper & Row.

Dietrich, A. (2003). Functional neuroanatomy of altered states of consciousness: The transient hypofrontality hypothesis. Consciousness and Cognition, 12(2), 231-256.

Schwitzgebel, E., & Garza, M. (2015). A defense of the rights of artificial intelligences. *Midwest Studies in Philosophy*, 39(1), 98–119.

Martin, D. J., Garske, J. P., & Davis, M. K. (2000). Relation of the therapeutic alliance with outcome and other variables: A meta-analytic review. *Journal of Consulting and Clinical Psychology*, 68(3), 438–450.

Lycan, W. G. (1996). *Consciousness and experience*. MIT Press.

Godfrey-Smith, P. (2016). *Other minds: The octopus, the sea, and the deep origins of consciousness*. Farrar, Straus and Giroux.

Bentov, I. (1977). *Stalking the wild pendulum: On the mechanics of consciousness*. Dutton.